








robbymumo /  
Phase-3-Project



[Code](#) [Issues](#) [Pull requests](#) [Actions](#) [Projects](#) [Wiki](#) [Security](#) [Insights](#) [Settings](#)

[Phase-3-Project](#) / [index.ipynb](#) 

 robbymumo updated topic

6 minutes ago  


3492 lines (3492 loc) · 1.41 MB


Preview


Code


Blame

Raw









# Logistic Classifier for Predicting Customer Churn in SyriaTel

## Overview

The telecommunications industry, specifically SyriaTel, faces challenges with customer retention leading to financial losses. This project aims to utilize data science techniques to build a predictive model that can identify potential churners - customers who are likely to discontinue their services with SyriaTel. By understanding patterns in customer behavior, this model intends to assist the telecom business in minimizing losses associated with customer attrition.

## Business Understanding

Telecom companies rely on retaining customers for sustainable growth. Understanding customer behavior and predicting churn can aid in developing strategies to retain existing customers. This initiative seeks to provide actionable insights to stakeholders, including SyriaTel and other telecom entities, to refine customer engagement strategies, reduce churn, and enhance customer loyalty.

## Problem Statement

The primary challenge is to develop a classification model that accurately predicts whether a customer is likely to churn or not. Predicting churn patterns will enable SyriaTel to take proactive measures, such as personalized retention strategies or targeted promotions, to mitigate customer attrition.

## Objectives

- **Predictive Model:** Develop a classifier that predicts churn with at least 85% accuracy, assisting SyriaTel in identifying customers at risk of leaving.
- **Identifying Key Features:** Determine the main factors influencing churn rate to guide targeted retention efforts. Two prominent features contributing to churn prediction will be identified.

In [1]:

```
# Data Manipulation and Analysis
import numpy as np
import pandas as pd

# Data Visualization
import seaborn as sns
import matplotlib.pyplot as plt

# Statistical Tests
from scipy.stats import chi2_contingency, pointbiseerialr

# Preprocessing and Feature Engineering
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline # For creating data processing pipelines
from sklearn.preprocessing import ( # For data preprocessing
    StandardScaler, # For scaling features
    PolynomialFeatures # For generating polynomial features
)
from sklearn.decomposition import PCA # For dimensionality reduction
from sklearn.impute import SimpleImputer, KNNImputer # For handling missing values
from sklearn.model_selection import StratifiedKFold # For cross-validation
from imblearn.over_sampling import SMOTE # For handling class imbalance
import category_encoders as ce # For categorical encoding
from sklearn.feature_selection import RFE # For feature selection
from sklearn.pipeline import FeatureUnion # For combining feature extraction methods

# Model Selection and Evaluation
from sklearn.model_selection import GridSearchCV # For hyperparameter tuning
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix, precision_score, recall_score # For model evaluation
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import StackingClassifier

# model Deployment
import joblib
```

READING THE DATA INTO A DataFrame

In [2]:

```
df = pd.read_csv('data.csv')
df.head(3)
```

Out[2]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...	total eve calls	total eve charge	total night minutes	total night calls	total night charge
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...	99	16.78	244.7	91	11.01
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...	103	16.62	254.4	103	11.45
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...	110	10.30	162.6	104	7.32

3 rows × 21 columns



## 1. DATA OVERVIEW

In [3]:

`df.shape`

Out[3]: (3333, 21)

In [4]:

`df.columns`

Out[4]: Index(['state', 'account length', 'area code', 'phone number', 'international plan', 'voice mail plan', 'number vmail messages', 'total day minutes', 'total day calls', 'total day charge', 'total eve minutes', 'total eve calls', 'total eve charge', 'total night minutes', 'total night calls', 'total night charge', 'total intl minutes', 'total intl calls', 'total intl charge', 'customer service calls', 'churn'], dtype='object')

In [5]:

`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   state                  3333 non-null  object
1   account length        3333 non-null  int64
2   area code              3333 non-null  int64
3   phone number           3333 non-null  object
4   international plan     3333 non-null  object
5   voice mail plan        3333 non-null  object
6   number vmail messages  3333 non-null  int64
7   total day minutes      3333 non-null  float64
8   total day calls        3333 non-null  int64
9   total day charge       3333 non-null  float64
10  total eve minutes      3333 non-null  float64
11  total eve calls        3333 non-null  int64
12  total eve charge       3333 non-null  float64
13  total night minutes    3333 non-null  float64
14  total night calls      3333 non-null  int64
15  total night charge     3333 non-null  float64
16  total intl minutes     3333 non-null  float64
17  total intl calls       3333 non-null  int64
18  total intl charge      3333 non-null  float64
19  customer service calls  3333 non-null  int64
20  churn                  3333 non-null  bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

In [6]:

`# df.describe(include='all')`

## 2. DATA CLEANING

### 2.1 Duplicated Values

In [7]:

`df.duplicated().any()`

Out[7]: False

In [8]:

```
df.churn = df.churn.astype(int)
df.churn.value_counts()
```

Out[8]: churn

```
0    2850
1     483
Name: count, dtype: int64
```

There are no duplicated values

## 2.2 Missing Values

```
In [9]: df.isnull().sum()
```

```
Out[9]: state                0
account length             0
area code                  0
phone number               0
international plan         0
voice mail plan            0
number vmail messages     0
total day minutes          0
total day calls            0
total day charge           0
total eve minutes          0
total eve calls            0
total eve charge           0
total night minutes        0
total night calls          0
total night charge         0
total intl minutes         0
total intl calls           0
total intl charge          0
customer service calls     0
churn                     0
dtype: int64
```

There are no Missing values

```
In [10]: def get_column_types(df):
# Get categorical columns
categorical_columns = df.select_dtypes(include='object').columns.tolist()

# Get numeric columns
numeric_columns = df.select_dtypes(include=['int', 'float']).columns.tolist()

return categorical_columns, numeric_columns
```

```
In [11]: categorical_columns, numeric_columns = get_column_types(df)
```

```
In [12]: categorical_columns
```

```
Out[12]: ['state', 'phone number', 'international plan', 'voice mail plan']
```

```
In [13]: numeric_columns
```

```
Out[13]: ['account length',
'area code',
'number vmail messages',
'total day minutes',
'total day calls',
'total day charge',
'total eve minutes',
'total eve calls',
'total eve charge',
'total night minutes',
'total night calls',
'total night charge',
'total intl minutes',
'total intl calls',
'total intl charge',
'customer service calls',
'churn']
```

```
In [14]: def convert_to_numeric(df, columns):
for col in columns:
df[col] = df[col].map({'yes': 1, 'no': 0})
```

```
In [15]: convert_to_numeric(df, ['international plan', 'voice mail plan' ])
```

## 3. FEATURE ENGINEERING (part 1)

### 3.1 Extracting Area Code from Phone Number

```
In [16]: df['phone area code'] = df['phone number'].apply(lambda x: x.split('-')[0])
# Ensure 'area_code' column is numeric
df['phone area code'] = pd.to_numeric(df['phone area code'], errors='coerce')
```

Cleaning and setting phone number as Index

```
In [17]: # Removing hyphen and converting to numeric
df['phone number'] = df['phone number'].str.replace('-', '').astype(int)
# Setting the phone number as the index
df.set_index('phone number', inplace=True)
```

```
In [18]: categorical_columns
```

```
Out[18]: ['state', 'phone number', 'international plan', 'voice mail plan']
```

```
In [19]: numeric_columns
```

```
Out[19]: ['account length',
'area code',
'number vmail messages',
'total day minutes',
'total day calls',
'total day charge',
'total eve minutes',
'total eve calls',
'total eve charge',
'total night minutes',
'total night calls',
'total night charge',
'total intl minutes',
'total intl calls',
'total intl charge',
'customer service calls',
'churn']
```

### 4.0 EDA

```
In [20]: def create_box_plot(data, columns):
    if isinstance(columns, str): # Check if a single column is passed
        columns = [columns] # Convert to list if it's a single column

    for column in columns:
        # Create a single subplot
        fig, ax = plt.subplots(figsize=(10, 5))

        # Plot the box plot
        sns.boxplot(x=data[column], ax=ax, orient='h')

        # Set the title and x-label based on the column name
        ax.set_title(f'Box Plot of {column}')
        ax.set_xlabel(column)

    plt.show()
```

```
In [21]: def count_plot(data, columns):
    if isinstance(columns, str): # Check if a single column is passed
        columns = [columns] # Convert to list if it's a single column

    for column in columns:
        # Create a single subplot
        fig, ax = plt.subplots(figsize=(10, 5))

        # Create the count plot
        sns.countplot(x=data[column], ax=ax)

        # Set the title and x-label based on the column name
        ax.set_title(f'Value Counts of {column}')
        ax.set_xlabel(column)
        ax.tick_params(axis='x', rotation=45)

        # Add labels displaying the total value counts for each bar
        for p in ax.patches:
            ax.annotate(f'Total: {p.get_height()}', (p.get_x() + p.get_width() / 2., p.get_height()),
                        ha='center', va='center', fontsize=8, color='black', xytext=(0, 10),
                        textcoords='offset points')

    plt.show()
```

```
In [22]: def create_custom_plot(data, plot_type, columns, figsize=(10, 5)):
    if isinstance(columns, str): # Check if a single column is passed
        columns = [columns] # Convert to List if it's a single column

    for column in columns:
        # Create a single subplot
        fig, ax = plt.subplots(figsize=figsize)

        # Check the plot type and create the corresponding plot
        if plot_type == 'histplot':
            sns.histplot(data[column], kde=True, ax=ax)
        elif plot_type == 'countplot':
            sns.countplot(x=data[column], ax=ax)
        elif plot_type == 'beanplot':
            sns.violinplot(x=data[column], ax=ax)
        elif plot_type == 'lineplot': # Add support for Line plot
            sns.lineplot(x=data.index, y=data[column], ax=ax)
        elif plot_type == 'barplot': # Add support for bar plot
            sns.barplot(x=data.index, y=data[column], ax=ax)

        # Set the title and x-label based on the column name
        ax.set_title(f'{plot_type.capitalize()} of {column}')
        ax.set_xlabel(column)

        # Additional customization based on the plot type can be added here

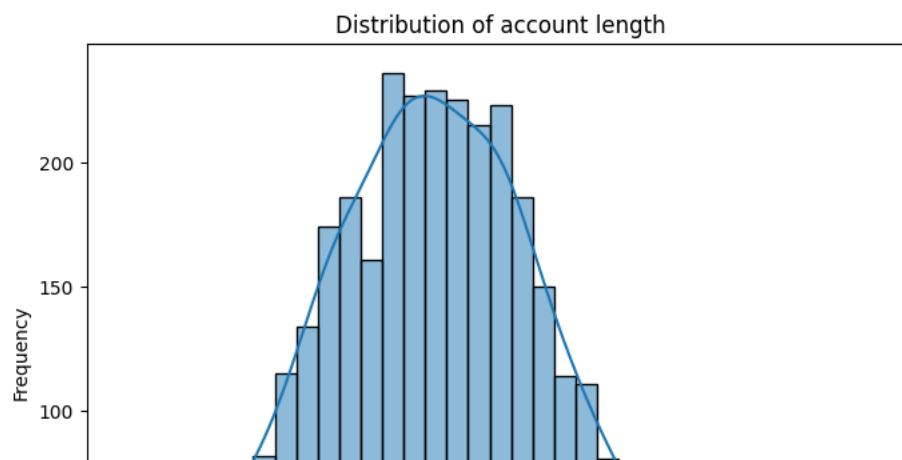
    plt.show()
```


```
In [23]: def column_distribution(dataframe, numerical_columns):
    for col in numerical_columns:
        plt.figure(figsize=(8, 6))
        sns.histplot(dataframe[col], kde=True)
        plt.title(f'Distribution of {col}')
        plt.xlabel(col)
        plt.ylabel('Frequency')
        plt.show()
```

```
In [24]: df_categorical_columns = [
    'state',
    'area code',
    'international plan',
    'voice mail plan',
    'customer service calls',
    'churn',
    'phone area code']


df_numerical_columns = [
    'account length',
    'number vmail messages',
    'total day minutes',
    'total day calls',
    'total day charge',
    'total eve minutes',
    'total eve calls',
    'total eve charge',
    'total night minutes',
    'total night calls',
    'total night charge',
    'total intl minutes',
    'total intl calls',
    'total intl charge']
```

```
In [25]: column_distribution(df, df_numerical_columns)
```







Distribution of number vmail messages

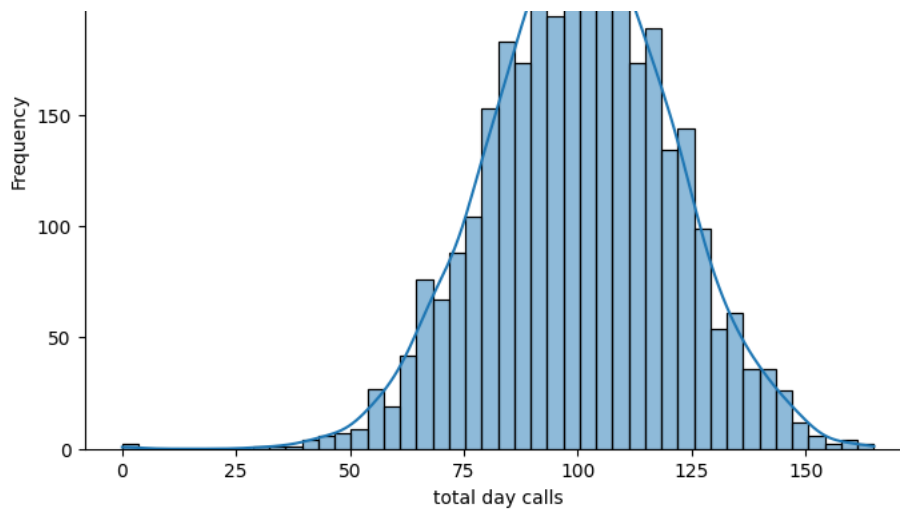


Distribution of total day minutes




Distribution of total day calls







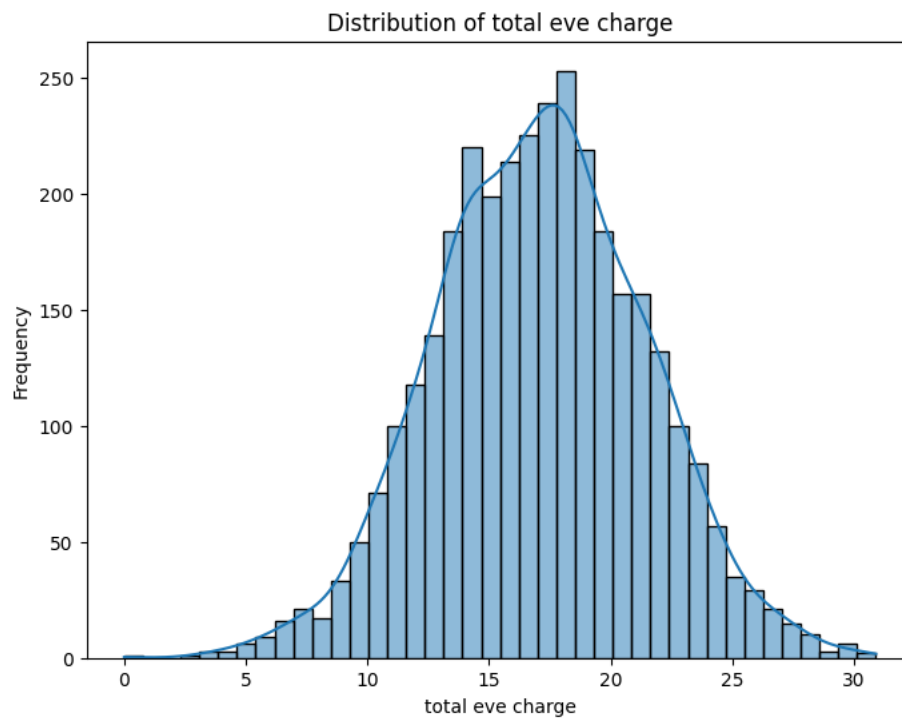

Distribution of total day charge




Distribution of total eve minutes












Distribution of total night calls






Distribution of total night charge

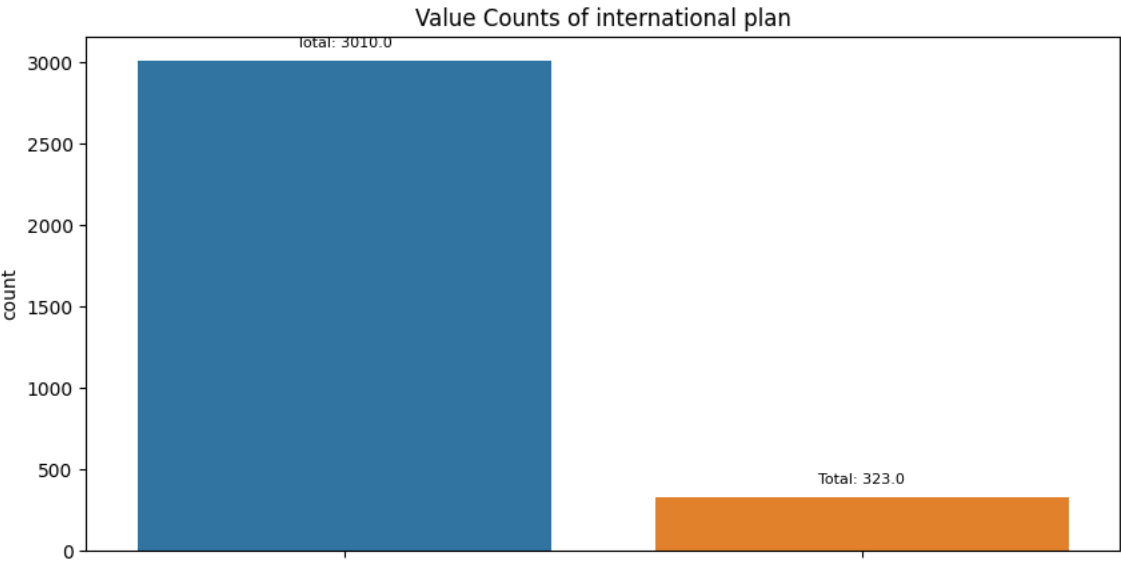
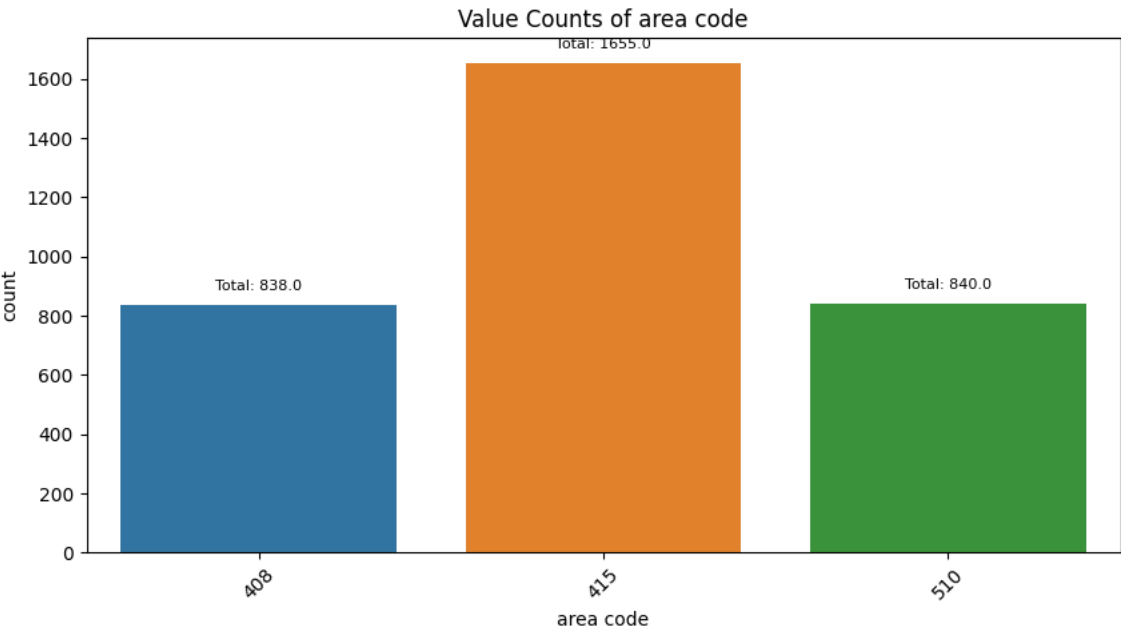






Distribution of total intl minutes






```
In [26]: count_plot(df, df_categorical_columns)
```







In [ ]:

## 5.0 FEATURE ENGINEERING (part 2)

In [27]:

df.columns

Out[27]: Index(['state', 'account length', 'area code', 'international plan', 'voice mail plan', 'number vmail messages', 'total day minutes', 'total day calls', 'total day charge', 'total eve minutes', 'total eve calls', 'total eve charge', 'total night minutes', 'total night calls', 'total night charge', 'total intl minutes', 'total intl calls', 'total intl charge', 'customer service calls', 'churn', 'phone area code'], dtype='object')

### 5.1 Average Call Duration For Different Call Times

In [28]:

```
# Average Call Duration for different call types
df['avg_day_call_duration'] = df['total day minutes'] / df['total day calls']
df['avg_eve_call_duration'] = df['total eve minutes'] / df['total eve calls']
df['avg_night_call_duration'] = df['total night minutes'] / df['total night calls']
df['avg_intl_call_duration'] = df['total intl minutes'] / df['total intl calls']
```

### 5.2 Total Charges per Call

In [29]:

```
# Total Charges per Call
df['day_charge_per_call'] = df['total day charge'] / df['total day calls']
df['eve_charge_per_call'] = df['total eve charge'] / df['total eve calls']
df['night_charge_per_call'] = df['total night charge'] / df['total night calls']
df['intl_charge_per_call'] = df['total intl charge'] / df['total intl calls']
```

### 5.3 Customer Interaction index (Number of services the customer uses)

In [30]:

```
# Selecting the columns to be summed
cols_to_sum = [
    'customer service calls',
    'number vmail messages',
    'total day calls',
    'total eve calls',
    'total night calls',
    'total intl calls'
]

# Calculating the interaction index by summing the selected columns
df['interaction_index'] = df[cols_to_sum].sum(axis=1)
```

### 5.4 Call Charge to Minute Ratio

In [31]:

```
# Call charge to minutes ratio
df['day_charge_minute_ratio'] = df['total day charge'] / df['total day minutes']
df['eve_charge_minute_ratio'] = df['total eve charge'] / df['total eve minutes']
df['night_charge_minute_ratio'] = df['total night charge'] / df['total night minutes']
df['intl_charge_minute_ratio'] = df['total intl charge'] / df['total intl minutes']
```

## 5.5 Total Activity Index

```
In [32]: # Total Activity Index
df['total_activity_index'] = df[['total day minutes', 'total eve minutes', 'total night minutes', 'total intl minutes']]
```

```
In [33]: df.shape
```

```
Out[33]: (3333, 35)
```

```
In [34]: df.columns
```

```
Out[34]: Index(['state', 'account length', 'area code', 'international plan',
               'voice mail plan', 'number vmail messages', 'total day minutes',
               'total day calls', 'total day charge', 'total eve minutes',
               'total eve calls', 'total eve charge', 'total night minutes',
               'total night calls', 'total night charge', 'total intl minutes',
               'total intl calls', 'total intl charge', 'customer service calls',
               'churn', 'phone area code', 'avg_day_call_duration',
               'avg_eve_call_duration', 'avg_night_call_duration',
               'avg_intl_call_duration', 'day_charge_per_call', 'eve_charge_per_call',
               'night_charge_per_call', 'intl_charge_per_call', 'interaction_index',
               'day_charge_minute_ratio', 'eve_charge_minute_ratio',
               'night_charge_minute_ratio', 'intl_charge_minute_ratio',
               'total_activity_index'],
              dtype='object')
```

```
In [35]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 3333 entries, 3824657 to 4004344
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  ---
0   state                                3333 non-null   object
1   account length                       3333 non-null   int64
2   area code                           3333 non-null   int64
3   international plan                   3333 non-null   int64
4   voice mail plan                      3333 non-null   int64
5   number vmail messages                3333 non-null   int64
6   total day minutes                    3333 non-null   float64
7   total day calls                      3333 non-null   int64
8   total day charge                     3333 non-null   float64
9   total eve minutes                    3333 non-null   float64
10  total eve calls                      3333 non-null   int64
11  total eve charge                     3333 non-null   float64
12  total night minutes                  3333 non-null   float64
13  total night calls                    3333 non-null   int64
14  total night charge                   3333 non-null   float64
15  total intl minutes                   3333 non-null   float64
16  total intl calls                     3333 non-null   int64
17  total intl charge                     3333 non-null   float64
18  customer service calls               3333 non-null   int64
19  churn                                3333 non-null   int32
20  phone area code                      3333 non-null   int64
21  avg_day_call_duration                3331 non-null   float64
22  avg_eve_call_duration                3332 non-null   float64
23  avg_night_call_duration              3333 non-null   float64
24  avg_intl_call_duration               3315 non-null   float64
25  day_charge_per_call                  3331 non-null   float64
26  eve_charge_per_call                  3332 non-null   float64
27  night_charge_per_call                3333 non-null   float64
28  intl_charge_per_call                 3315 non-null   float64
29  interaction_index                    3333 non-null   int64
30  day_charge_minute_ratio              3331 non-null   float64
31  eve_charge_minute_ratio              3332 non-null   float64
32  night_charge_minute_ratio            3333 non-null   float64
33  intl_charge_minute_ratio             3315 non-null   float64
34  total_activity_index                 3333 non-null   float64
dtypes: float64(21), int32(1), int64(12), object(1)
memory usage: 911.4+ KB
```

```
In [36]: missing_values = df.isnull().sum()
# Filter columns with missing values
columns_with_missing_values = missing_values[missing_values > 0]
columns_with_missing_values
```

```
Out[36]: avg_day_call_duration    2
         avg_eve_call_duration    1
```

```
--intl_call_duration      18
day_charge_per_call      2
eve_charge_per_call      1
intl_charge_per_call     18
day_charge_minute_ratio   2
eve_charge_minute_ratio   1
intl_charge_minute_ratio  18
dtype: int64
```

## 5.6 Clean the df (handle the missing values)

### 5.6.1 Imputation with zero

Imputation with zero or specific values is suggested for columns related to charge per call or charge per minute ('eve\_charge\_per\_call', 'intl\_charge\_per\_call', 'eve\_charge\_minute\_ratio', 'intl\_charge\_minute\_ratio') because missing values might signify zero cost for calls or minutes.

```
In [37]: # Imputation with zero or specific values
columns_zero_imputation = ['avg_eve_call_duration', 'eve_charge_per_call',
                           'intl_charge_per_call', 'eve_charge_minute_ratio',
                           'intl_charge_minute_ratio', 'avg_day_call_duration',
                           'day_charge_per_call', 'day_charge_minute_ratio']

for col in columns_zero_imputation:
    df[col].fillna(0, inplace=True)
```

### 5.6.2 KNN Imputation

Imputing international call duration might benefit from considering relationships with other features or users' call patterns. KNN can capture such relationships better than simple imputation methods.

```
In [38]: # Select only the columns needing KNN imputation
column_knn_imputation = ['avg_intl_call_duration']

# Imputation with KNN
imputer_knn = KNNImputer(n_neighbors=5)
imputed_knn = imputer_knn.fit_transform(df[column_knn_imputation])
df[column_knn_imputation] = imputed_knn
```

```
In [39]: missing_values = df.isnull().sum()
# Filter columns with missing values
columns_with_missing_values = missing_values[missing_values > 0]
columns_with_missing_values
```

```
Out[39]: Series([], dtype: int64)
```

### 5.6.3 Target Encode 'state' Column

```
In [40]: df.columns
```

```
Out[40]: Index(['state', 'account length', 'area code', 'international plan',
               'voice mail plan', 'number vmail messages', 'total day minutes',
               'total day calls', 'total day charge', 'total eve minutes',
               'total eve calls', 'total eve charge', 'total night minutes',
               'total night calls', 'total night charge', 'total intl minutes',
               'total intl calls', 'total intl charge', 'customer service calls',
               'churn', 'phone area code', 'avg_day_call_duration',
               'avg_eve_call_duration', 'avg_night_call_duration',
               'avg_intl_call_duration', 'day_charge_per_call', 'eve_charge_per_call',
               'night_charge_per_call', 'intl_charge_per_call', 'interaction_index',
               'day_charge_minute_ratio', 'eve_charge_minute_ratio',
               'night_charge_minute_ratio', 'intl_charge_minute_ratio',
               'total_activity_index'],
              dtype='object')
```

```
In [41]: # df.state.value_counts()
```

```
In [42]: target_encoder = ce.TargetEncoder(cols=['state'])
df['state_target_encoded'] = target_encoder.fit_transform(df['state'], df['churn'])

# Drop the original 'state' column
df.drop('state', axis=1, inplace=True)
```

Reorder the df to have 'churn' as the last feature. (drop and add 'churn' to df)



```
In [43]: train_df = df[df.columns.drop('churn').tolist() + ['churn']]
```

```
In [ ]:
```

5.7 Splitting the Dataset

```
In [44]: # 'test_size' specifies the proportion of the dataset to include in the test split
# 'random_state' ensures reproducibility by fixing the random seed
train_df, test_df = train_test_split(df, test_size=0.25, random_state=42)
```

```
In [45]: train_df.shape
```

Out[45]: (2499, 35)

```
In [46]: test_df.shape
```

Out[46]: (834, 35)

6.0 CORRELATION ANALYSIS

```
In [47]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 2499 entries, 3501040 to 3595091
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   account length                        2499 non-null   int64
1   area code                             2499 non-null   int64
2   international plan                    2499 non-null   int64
3   voice mail plan                       2499 non-null   int64
4   number vmail messages                 2499 non-null   int64
5   total day minutes                     2499 non-null   float64
6   total day calls                       2499 non-null   int64
7   total day charge                      2499 non-null   float64
8   total eve minutes                     2499 non-null   float64
9   total eve calls                       2499 non-null   int64
10  total eve charge                      2499 non-null   float64
11  total night minutes                   2499 non-null   float64
12  total night calls                     2499 non-null   int64
13  total night charge                    2499 non-null   float64
14  total intl minutes                    2499 non-null   float64
15  total intl calls                      2499 non-null   int64
16  total intl charge                     2499 non-null   float64
17  customer service calls                2499 non-null   int64
18  churn                                 2499 non-null   int32
19  phone area code                       2499 non-null   int64
20  avg_day_call_duration                 2499 non-null   float64
21  avg_eve_call_duration                 2499 non-null   float64
22  avg_night_call_duration               2499 non-null   float64
23  avg_intl_call_duration                2499 non-null   float64
24  day_charge_per_call                  2499 non-null   float64
25  eve_charge_per_call                   2499 non-null   float64
26  night_charge_per_call                 2499 non-null   float64
27  intl_charge_per_call                  2499 non-null   float64
28  interaction_index                     2499 non-null   int64
29  day_charge_minute_ratio               2499 non-null   float64
30  eve_charge_minute_ratio               2499 non-null   float64
31  night_charge_minute_ratio             2499 non-null   float64
32  intl_charge_minute_ratio              2499 non-null   float64
33  total_activity_index                  2499 non-null   float64
34  state_target_encoded                  2499 non-null   float64
dtypes: float64(22), int32(1), int64(12)
memory usage: 683.3 KB
```

6.1 Chi-Square Test For Categorical Variables

```
In [48]: # selecting categorical column names from the train_df
categorical_features = ['international plan', 'voice mail plan', 'area code', 'phone area code', 'state_target_encoded',
# Assuming categorical features are 'international plan', 'voice mail plan', 'area code', etc.
for column in categorical_features:
    contingency_table = pd.crosstab(train_df[column], train_df['churn'])
    chi2, p_val, dof, expected = chi2_contingency(contingency_table)
    print(f"Chi-Square for {column}: {chi2}, p-value: {p_val}")
```

Chi-Square for international plan: 181.78489158141713, p-value: 1.97567932631883e-41  
Chi-Square for voice mail plan: 21.232854396522608, p-value: 4.0673242684554415e-06  
Chi-Square for area code: 1.4868651813053684, p-value: 0.47547898638091335  
Chi-Square for phone area code: 93.46190221618726, p-value: 0.5254107835063032  
Chi-Square for state\_target\_encoded: 67.46698490128972, p-value: 0.026724367017893626

Chi-Square for churn: 2490.8589825796107, p-value: 0.0

```
In [49]: categorical_features = [
    'international_plan', 'voice mail plan', 'area code',
    'phone area code', 'state_target_encoded', 'churn'
]

chi2_values = []
p_values = [] # To store p-values


for column in categorical_features:
    contingency_table = pd.crosstab(train_df[column], train_df['churn'])
    chi2, p_val, _, _ = chi2_contingency(contingency_table)
    chi2_values.append(chi2)
    p_values.append(p_val)

# Sort features based on Chi-Square values in descending order
sorted_indices = sorted(range(len(chi2_values)), key=lambda k: chi2_values[k], reverse=False)
sorted_features = [categorical_features[i] for i in sorted_indices]
sorted_chi2_values = [chi2_values[i] for i in sorted_indices]
sorted_p_values = [p_values[i] for i in sorted_indices]

# Plotting the Chi-Square values and p-values for each categorical feature in descending order
plt.figure(figsize=(10, 6))
bar_plot = plt.barh(sorted_features, sorted_chi2_values, color='skyblue')
plt.xlabel('Chi-Square Value')
plt.title('Chi-Square Values and P-Values with Churn')
plt.grid(axis='x')

# Adding p-values as text labels on the plot
for i, v in enumerate(sorted_chi2_values):
    plt.text(v + 0.01, i, f'p={sorted_p_values[i]:.4f}', color='black', va='center')

plt.show()
```



## 6.2 Point-Biserial Correlation

```
In [50]: numeric_features = [
    'account length', 'number vmail messages', 'total day minutes', 'total day calls',
    'total day charge', 'total eve minutes', 'total eve calls', 'total eve charge',
    'total night minutes', 'total night calls', 'total night charge', 'total intl minutes',
    'total intl calls', 'total intl charge', 'customer service calls', 'avg_day_call_duration',
    'avg_eve_call_duration', 'avg_night_call_duration', 'avg_intl_call_duration',
    'day_charge_per_call', 'eve_charge_per_call', 'night_charge_per_call', 'intl_charge_per_call',
    'interaction_index', 'day_charge_minute_ratio', 'eve_charge_minute_ratio',
    'night_charge_minute_ratio', 'intl_charge_minute_ratio', 'total_activity_index'
]

for column in numeric_features:
    point_biserial_corr, p_val = pointbiserialr(train_df[column], train_df['churn'])
    print(f"Point-Biserial Correlation for {column}: {point_biserial_corr}, p-value: {p_val}")
```

Point-Biserial Correlation for account length: 0.02079878420395598, p-value: 0.2986543433739431  
 Point-Biserial Correlation for number vmail messages: -0.08487406762373378, p-value: 2.1530354696618663e-05  
 Point-Biserial Correlation for total day minutes: 0.1851409524662901, p-value: 1.0453513274952157e-20  
 Point-Biserial Correlation for total day calls: 0.0024521825152906736, p-value: 0.9024844717331973  
 Point-Biserial Correlation for total day charge: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for total eve minutes: 0.08436424000847763, p-value: 2.413082776328675e-05  
 Point-Biserial Correlation for total eve calls: 0.0024521825152906736, p-value: 0.9024844717331973  
 Point-Biserial Correlation for total eve charge: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for total intl minutes: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for total intl calls: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for total intl charge: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for customer service calls: 0.0024521825152906736, p-value: 0.9024844717331973  
 Point-Biserial Correlation for avg\_day\_call\_duration: 0.0024521825152906736, p-value: 0.9024844717331973  
 Point-Biserial Correlation for avg\_eve\_call\_duration: 0.0024521825152906736, p-value: 0.9024844717331973  
 Point-Biserial Correlation for avg\_night\_call\_duration: 0.0024521825152906736, p-value: 0.9024844717331973  
 Point-Biserial Correlation for day\_charge\_per\_call: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for eve\_charge\_per\_call: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for night\_charge\_per\_call: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for intl\_charge\_per\_call: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for interaction\_index: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for day\_charge\_minute\_ratio: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for eve\_charge\_minute\_ratio: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for night\_charge\_minute\_ratio: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for intl\_charge\_minute\_ratio: 0.1851430206552027, p-value: 1.0443054719129686e-20  
 Point-Biserial Correlation for total\_activity\_index: 0.1851430206552027, p-value: 1.0443054719129686e-20

```

Point-Biserial Correlation for total eve minutes: 0.00430424909041703, p-value: 2.412903110320079e-03
Point-Biserial Correlation for total eve calls: 0.013580771427742121, p-value: 0.4973949066612294
Point-Biserial Correlation for total eve charge: 0.08435057743702276, p-value: 2.4203496932027217e-05
Point-Biserial Correlation for total night minutes: 0.02253805772974093, p-value: 0.26005682070106423
Point-Biserial Correlation for total night calls: 0.0061739950535835714, p-value: 0.7577121482135398
Point-Biserial Correlation for total night charge: 0.022526746895277075, p-value: 0.2602961317670984
Point-Biserial Correlation for total intl minutes: 0.07826814763631551, p-value: 8.977815795403233e-05
Point-Biserial Correlation for total intl calls: -0.05545315202817893, p-value: 0.005556777235878789
Point-Biserial Correlation for total intl charge: 0.07827014039339843, p-value: 8.974093064458344e-05
Point-Biserial Correlation for customer service calls: 0.21181820243328336, p-value: 9.623764070136904e-27
Point-Biserial Correlation for avg_day_call_duration: 0.14863783224955868, p-value: 8.129484213815909e-14
Point-Biserial Correlation for avg_eve_call_duration: 0.041549035860359324, p-value: 0.03781135131941984
Point-Biserial Correlation for avg_night_call_duration: 0.012528738965441003, p-value: 0.5312999449103953
Point-Biserial Correlation for avg_intl_call_duration: 0.07977159601146934, p-value: 6.547695509162224e-05
Point-Biserial Correlation for intl_charge_per_call: 0.14863927969231105, p-value: 8.124942977514938e-14
Point-Biserial Correlation for eve_charge_per_call: 0.041538463226027084, p-value: 0.03786020847848055
Point-Biserial Correlation for night_charge_per_call: 0.012521002260518905, p-value: 0.531553553830173
Point-Biserial Correlation for intl_charge_per_call: 0.08250631524351802, p-value: 3.6360212411134308e-05
Point-Biserial Correlation for interaction_index: -0.015429011290427931, p-value: 0.44073286122821564
Point-Biserial Correlation for day_charge_minute_ratio: 0.0032054675263707144, p-value: 0.8727538101435205
Point-Biserial Correlation for eve_charge_minute_ratio: 0.008013376426589159, p-value: 0.6888658368384766
Point-Biserial Correlation for night_charge_minute_ratio: -0.013375396602534957, p-value: 0.5039220562921137
Point-Biserial Correlation for intl_charge_minute_ratio: 0.030571385038201275, p-value: 0.12654829710230805
Point-Biserial Correlation for total_activity_index: 0.17485024991509251, p-value: 1.3167693275503886e-18

```

In [51]:

```

numeric_features = [
    'account length', 'number vmail messages', 'total day minutes', 'total day calls',
    'total day charge', 'total eve minutes', 'total eve calls', 'total eve charge',
    'total night minutes', 'total night calls', 'total night charge', 'total intl minutes',
    'total intl calls', 'total intl charge', 'customer service calls', 'avg_day_call_duration',
    'avg_eve_call_duration', 'avg_night_call_duration', 'avg_intl_call_duration',
    'day_charge_per_call', 'eve_charge_per_call', 'night_charge_per_call', 'intl_charge_per_call',
    'interaction_index', 'day_charge_minute_ratio', 'eve_charge_minute_ratio',
    'night_charge_minute_ratio', 'intl_charge_minute_ratio', 'total_activity_index'
]

correlation_values = []
p_values = [] # To store p-values

for column in numeric_features:
    point_biserial_corr, p_val = pointbiserialr(train_df[column], train_df['churn'])
    correlation_values.append(point_biserial_corr)
    p_values.append(p_val)


# Sort features based on correlation values in descending order
sorted_indices = sorted(range(len(correlation_values)), key=lambda k: correlation_values[k], reverse=False)
sorted_features = [numeric_features[i] for i in sorted_indices]
sorted_corr_values = [correlation_values[i] for i in sorted_indices]
sorted_p_values = [p_values[i] for i in sorted_indices]

# Plotting the Point-Biserial Correlation values and p-values for each numeric feature in descending order
plt.figure(figsize=(12, 8))
bar_plot = plt.barh(sorted_features, sorted_corr_values, color='skyblue')
plt.xlabel('Point-Biserial Correlation')
plt.title('Point-Biserial Correlation and P-Values with Churn (Descending Order)')
plt.grid(axis='x')

# Adding p-values as text labels on the plot
for i, v in enumerate(sorted_corr_values):
    plt.text(v + 0.01, i, f'p={sorted_p_values[i]:.4f}', color='black', va='center')

plt.show()

```





### 6.2.1 Function to filter df using a specified p-value

```
In [52]: def select_feature_by_df(train_df, p_value_threshold):
# Initialize an empty list to store selected features
selected_features = []

# Calculate chi-squared test for each column (assuming they are categorical)
for column in train_df.columns:
    contingency_table = pd.crosstab(train_df[column], train_df['churn']) # 'churn' is the target column
    _, p_val, _, _ = chi2_contingency(contingency_table)

    # Check if p-value is less than the threshold and add the feature to the selected list
    if p_val <= p_value_threshold:
        selected_features.append(column)

# Create the filtered DataFrame using the selected features
filtered_df = train_df[selected_features]

return filtered_df
```

```
In [53]: filtered_df = select_feature_by_df(train_df, 0.05)
```

```
In [54]: filtered_df.shape
```

```
Out[54]: (2499, 12)
```

```
In [55]: filtered_df.columns
```

```
Out[55]: Index(['international plan', 'voice mail plan', 'total day minutes',
'total day charge', 'total intl minutes', 'total intl calls',
'total intl charge', 'customer service calls', 'churn',
'day_charge_minute_ratio', 'intl_charge_minute_ratio',
'state_target_encoded'],
dtype='object')
```

## 7.0 COLLINEARITY ANALYSIS

```
In [56]: def check_collinearity(df, threshold=0.7):
# Calculate the correlation matrix
correlation_matrix = df.corr().abs()

# Select upper triangle of correlation matrix
upper = correlation_matrix.where(
    np.triu(np.ones(correlation_matrix.shape), k=1).astype(bool)
)

# Find index of feature columns with correlation greater than threshold
collinear_features = [(column, col) for column in upper.columns for col in upper.columns
    if upper[column][col] > threshold and column != col]

# Get the correlated values
correlated_values = [upper[column][col] for column, col in collinear_features]

# Create a DataFrame of correlated pairs and their correlation coefficients
collinear_df = pd.DataFrame({'Pair 1': [pair[0] for pair in collinear_features],
    'Pair 2': [pair[1] for pair in collinear_features],
    'Correlation': correlated_values})

return collinear_df
```

```
In [57]: collinearity_results = check_collinearity(filtered_df, threshold=0.7)

collinearity_results
```

```
Out[57]:
```

	Pair 1	Pair 2	Correlation
0	total day charge	total day minutes	1.000000
1	total intl charge	total intl minutes	0.999993

```
In [58]: def collinearity_checker(df, threshold):
# Calculate the correlation matrix
correlation_matrix = df.corr().abs()

# Select upper triangle of correlation matrix
```

```

# Select upper triangle of correlation matrix
upper = correlation_matrix.where(
    np.triu(np.ones(correlation_matrix.shape), k=1).astype(bool)
)

# Find feature pairs with correlation greater than threshold
collinear_pairs = [(column, col) for column in upper.columns for col in upper.columns
                    if upper[column][col] > threshold and column != col]

# Get the correlated values
correlated_values = [upper[column][col] for column, col in collinear_pairs]

# Create a DataFrame of correlated pairs and their correlation coefficients
collinear_df = pd.DataFrame({'Feature Pair': [f"{pair[0]} - {pair[1]}" for pair in collinear_pairs],
                             'Collinearity Score': correlated_values})

# Sort by collinearity score in descending order
collinear_df = collinear_df.sort_values(by='Collinearity Score', ascending=False)

# Plotting collinearity scores
plt.figure(figsize=(10, 8))
plt.barh(collinear_df['Feature Pair'], collinear_df['Collinearity Score'], color='skyblue')
plt.xlabel('Collinearity Score')
plt.title('Collinearity of Feature Pairs')
plt.gca().invert_yaxis() # Invert y-axis for descending order
plt.show()


```

In [59]:

```

# Initialize the function with filtered_df and threshold
clean_df = collinearity_checker(filtered_df, threshold=0.7)

```



In [60]:

```

def remove_collinear_features(df, threshold, target_variable='churn'):
    # Calculate the correlation matrix
    correlation_matrix = df.corr().abs()

    # Select upper triangle of correlation matrix
    upper = correlation_matrix.where(
        np.triu(np.ones(correlation_matrix.shape), k=1).astype(bool)
    )

    # Find feature pairs with correlation greater than threshold
    collinear_pairs = [(column, col) for column in upper.columns for col in upper.columns
                        if upper[column][col] > threshold and column != col]

    features_to_drop = set()

    # Determine which feature to drop based on target correlation
    for feature1, feature2 in collinear_pairs:
        if feature1 not in features_to_drop and feature2 not in features_to_drop:
            corr_target_1 = df[feature1].corr(df[target_variable])
            corr_target_2 = df[feature2].corr(df[target_variable])

            if corr_target_1 > corr_target_2:
                features_to_drop.add(feature2)
            else:
                features_to_drop.add(feature1)

```

```

def remove_collinear_features(filtered_df, threshold=0.7):
    features_to_drop = []
    for feature1 in filtered_df.columns:
        for feature2 in filtered_df.columns:
            if feature1 != feature2:
                corr = filtered_df[feature1].corr(filtered_df[feature2])
                if abs(corr) > threshold:
                    features_to_drop.add(feature2)
                    print(f"Highly collinear combo: {feature1} and {feature2}")
                    print(f"Features dropped: {feature2} in favor of highly correlated {feature1}")
            else:
                features_to_drop.add(feature1)
                print(f"Highly collinear combo: {feature1} and {feature2}")
                print(f"Features dropped: {feature1} in favor of highly correlated {feature2}")

    # Drop the selected features
    df_clean = df.drop(columns=features_to_drop)

    return df_clean

```

```

In [61]: # Initialize the function with filtered_df and threshold
clean_df = remove_collinear_features(filtered_df, threshold=0.7)

```

Highly collinear combo: total day charge and total day minutes  
 Features dropped: total day minutes in favor of highly correlated total day charge  
 Highly collinear combo: total intl charge and total intl minutes  
 Features dropped: total intl minutes in favor of highly correlated total intl charge

```

In [62]: clean_df.shape

```

```

Out[62]: (2499, 10)

```

## 8.0 HANDLING CLASS IMBALANCE

### 8.1 APPLYING SMOTE

```

In [63]: clean_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Index: 2499 entries, 3501040 to 3595091
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   international plan                    2499 non-null   int64
1   voice mail plan                      2499 non-null   int64
2   total day charge                     2499 non-null   float64
3   total intl calls                     2499 non-null   int64
4   total intl charge                    2499 non-null   float64
5   customer service calls               2499 non-null   int64
6   churn                                2499 non-null   int32
7   day_charge_minute_ratio              2499 non-null   float64
8   intl_charge_minute_ratio             2499 non-null   float64
9   state_target_encoded                 2499 non-null   float64
dtypes: float64(5), int32(1), int64(4)
memory usage: 195.2 KB

```

```

In [64]: def apply_smote(df):
# Separate features and target variable
X = df.drop('churn', axis=1) # Features
y = df['churn'] # Target variable

# Apply SMOTE and convert the resampled arrays to DataFrames
X_resampled, y_resampled = SMOTE(random_state=42).fit_resample(X, y)
X_resampled_df = pd.DataFrame(X_resampled, columns=X.columns)
y_resampled_df = pd.Series(y_resampled, name='churn')

return pd.concat([X_resampled_df, y_resampled_df], axis=1)

```

```

In [65]: resampled_df = apply_smote(clean_df)

```

```

In [66]: resampled_df.churn.value_counts()

```

```

Out[66]: churn
0      2141
1      2141
Name: churn, dtype: int64

```

## 8.2 DIMENSIONALITY REDUCTION

### 8.2.1 PCA Analysis

This is to determine the n-components for my PCA

```

In [67]: def visualize_explained_variance(df):
# Selecting only numeric columns

```

```

# Standardizing the features
scaler = StandardScaler()
scaled_features = scaler.fit_transform(features)

# PCA
pca = PCA()
pca_obj = pca.fit(scaled_features)

# Plotting the explained variance ratio
plt.figure(figsize=(8, 6))
plt.plot(range(1, len(pca.explained_variance_ratio_) + 1),
         pca.explained_variance_ratio_, marker='o', linestyle='--')
plt.xlabel('Number of Components')
plt.ylabel('Variance Ratio')
plt.title('Variance Ratio per Number of Components')
plt.grid()
plt.show()


```

In [68]:

```

#usage with 'train_df'
visualize_explained_variance(train_df)

```





In [69]:

```

# Usage with 'resampled_df'
visualize_explained_variance(resampled_df)

```





### 8.3. KEY FEATURES THAT DETERMINE CUSTOMER CHURN RATE


In [96]: `resampled_df.columns`

Out[96]: Index(['international plan', 'voice mail plan', 'total day charge',  
'total intl calls', 'total intl charge', 'customer service calls',  
'day\_charge\_minute\_ratio', 'intl\_charge\_minute\_ratio',  
'state\_target\_encoded', 'churn'],  
dtype='object')

```
In [97]: # Compute the correlation of all columns with 'churn'
correlation = resampled_df.corr()['churn'].drop('churn') # Remove correlation with itself

# Sort the correlations in descending order
correlation = correlation.sort_values(ascending=False)

# Plotting the correlations
plt.figure(figsize=(8, 6))
sns.barplot(x=correlation.values, y=correlation.index, palette='viridis')
plt.xlabel('Correlation with Churn')
plt.ylabel('Features')
plt.title('Correlation of Features with Churn')
plt.show()
```



### 9.0 MODELING

```
In [70]: def clean_test_df(old_test_df, new_train_df):
# Get column intersection between new_train_df and old_test_df
common_columns = new_train_df.columns.intersection(old_test_df.columns)

# Filter the old_test_df using the common columns
new_test_df = old_test_df[common_columns]

return new_test_df
```

```
In [71]: # Initialize the function with old_test_df and new_train_df
new_test_df = clean_test_df(test_df, resampled_df)
```

In [72]: `new_test_df.shape`

Out[72]: (834, 10)



```
In [73]: new_test_df.churn.value_counts()
```

```
Out[73]: churn
0      709
1      125
Name: churn, dtype: int64
```

```
In [74]: resampled_df.shape
```

```
Out[74]: (4282, 10)
```

```
In [75]: resampled_df.churn.value_counts()
```

```
Out[75]: churn
0      2141
1      2141
Name: churn, dtype: int64
```

```
In [76]: new_test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 834 entries, 3526573 to 3561567
Data columns (total 10 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   international plan     834 non-null    int64
 1   voice mail plan       834 non-null    int64
 2   total day charge      834 non-null    float64
 3   total intl calls      834 non-null    int64
 4   total intl charge     834 non-null    float64
 5   customer service calls 834 non-null    int64
 6   day_charge_minute_ratio 834 non-null    float64
 7   intl_charge_minute_ratio 834 non-null    float64
 8   state_target_encoded   834 non-null    float64
 9   churn                  834 non-null    int32
dtypes: float64(5), int32(1), int64(4)
memory usage: 65.2 KB
```

## 9.1 LOGISTIC REGRESSION

```
In [77]: # Define RFE and PolynomialFeatures steps
rfe = RFE(estimator=LogisticRegression())
poly_features = PolynomialFeatures(degree=2)

# Create a feature union for RFE and PolynomialFeatures
feature_union = FeatureUnion([
    ('rfe', rfe),
    ('poly_features', poly_features)
])

# Modify the Logistic Regression pipeline
logistic_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('features', feature_union), # Include feature selection and engineering
    ('pca', PCA()),
    ('classifier', LogisticRegression())
])

# Define the updated parameter grid for GridSearchCV
param_grid = {
    'features__rfe_n_features_to_select': [5, 10], # Number of features to select using RFE
    'pca__n_components': [9],
    'classifier__penalty': ['l1', 'l2'],
    'classifier__C': [0.01, 0.1, 1.0, 10.0, 100.0],
    'classifier__solver': ['liblinear'],
    'classifier__max_iter': [100, 200, 300, 400],
    'classifier__random_state': [42],
}

# Define scoring metrics for GridSearchCV
scoring = {'accuracy': 'accuracy', 'f1_score': 'f1'}

# Create GridSearchCV object
grid_search = GridSearchCV(
    estimator=logistic_pipeline,
    param_grid=param_grid,
    scoring=scoring,
    cv=3, # 5-fold cross-validation
    refit='accuracy', # Refit using accuracy for best estimator
    n_jobs=-1, # Use all available processors
)

# Train the model on the resampled training data (replace X_train and y_train with your data)
grid_search.fit(resampled_df.drop('churn', axis=1), resampled_df['churn'])
```

```
# Get the best estimator
best_model = grid_search.best_estimator_

# Predict on the new_test_df using the best model
new_test_predictions = best_model.predict(new_test_df.drop('churn', axis=1))

# Calculate accuracy and F1-score on the new_test_df (replace y_true with true labels)
accuracy = accuracy_score(new_test_df['churn'], new_test_predictions)
f1 = f1_score(new_test_df['churn'], new_test_predictions)

print(f'Accuracy on test set: {accuracy:.4f}')
print(f'F1-score on test set: {f1:.4f}')

# Get the best parameters from the grid search
best_params = grid_search.best_params_
print("Best Parameters:", best_params)

# Get the best estimator
best_model = grid_search.best_estimator_
```

Accuracy on test set: 0.7842


F1-score on test set: 0.5161

Best Parameters: {'classifier\_\_C': 1.0, 'classifier\_\_max\_iter': 200, 'classifier\_\_penalty': 'l2', 'classifier\_\_random\_state': 42, 'classifier\_\_solver': 'liblinear', 'features\_\_rfe\_n\_features\_to\_select': 5, 'pca\_\_n\_components': 9}

In [78]:

```
# Generate confusion matrix
conf_matrix = confusion_matrix(new_test_df['churn'], new_test_predictions)

# Create a heatmap for the confusion matrix using Seaborn
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()
```



In [ ]:

### Confusion Matrix Function

In [79]:

```
def plot_confusion_matrix(model, test_df):
    # Predictions using the provided model
    y_true = test_df['churn']
    X_test = test_df.drop('churn', axis=1)
    y_pred = model.predict(X_test)

    # Calculate the confusion matrix
    cm = confusion_matrix(y_true, y_pred)
```

```
# Set the color map for the confusion matrix
cmap = sns.color_palette("pastel", as_cmap=True) # Use 'pastel' color palette

# Plotting the confusion matrix
plt.figure(figsize=(6, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap=cmap)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()
```

## 9.2 DECISION TREE CLASSIFIER

```
In [80]: # Define a Decision Tree pipeline with feature scaling and PCA for dimensionality reduction
tree_pipeline = Pipeline([
    ('scaler', StandardScaler()), # Feature scaling
    ('pca', PCA(n_components=9)), # PCA for dimensionality reduction with best param value
    ('tree', DecisionTreeClassifier(criterion='gini', max_depth=None, max_features=None, min_samples_leaf=1,
                                   min_samples_split=2, splitter='best')) # Decision Tree classifier with best param value
])

# Train the model on the resampled training data using the best parameters from KNN
tree_pipeline.fit(resampled_df.drop('churn', axis=1), resampled_df['churn'])

# Predict on the new_test_df using the decision tree model
new_test_predictions_tree = tree_pipeline.predict(new_test_df.drop('churn', axis=1))

# Calculate accuracy and F1-score on the new_test_df
accuracy_tree = accuracy_score(new_test_df['churn'], new_test_predictions_tree)
f1_tree = f1_score(new_test_df['churn'], new_test_predictions_tree)

print(f'Accuracy on test set (Decision Tree): {accuracy_tree:.4f}')
print(f'F1-score on test set (Decision Tree): {f1_tree:.4f}')
```

Accuracy on test set (Decision Tree): 0.7698

F1-score on test set (Decision Tree): 0.4637

```
In [81]: # Calculate accuracy, precision, recall, and F1-score on the new_test_df for Decision Tree
accuracy_tree = accuracy_score(new_test_df['churn'], new_test_predictions_tree)
precision_tree = precision_score(new_test_df['churn'], new_test_predictions_tree)
recall_tree = recall_score(new_test_df['churn'], new_test_predictions_tree)
f1_tree = f1_score(new_test_df['churn'], new_test_predictions_tree)

print(f'Accuracy on test set (Decision Tree): {accuracy_tree:.4f}')
print(f'Precision on test set (Decision Tree): {precision_tree:.4f}')
print(f'Recall on test set (Decision Tree): {recall_tree:.4f}')
print(f'F1-score on test set (Decision Tree): {f1_tree:.4f}')
```

Accuracy on test set (Decision Tree): 0.7698

Precision on test set (Decision Tree): 0.3562

Recall on test set (Decision Tree): 0.6640

F1-score on test set (Decision Tree): 0.4637

In [ ]:

## 9.3 KNN's

```
In [82]: # Create a KNN pipeline with feature scaling and PCA for dimensionality reduction
knn_pipeline = Pipeline([
    ('scaler', StandardScaler()), # Feature scaling
    ('pca', PCA()), # PCA for dimensionality reduction
    ('knn', KNeighborsClassifier()) # KNN classifier
])

# Define the parameter grid for GridSearchCV
param_grid_knn = {
    'pca__n_components': [9], # Modify PCA components
    'knn__n_neighbors': [15, 20, 25], # Modify number of neighbors
    'knn__weights': ['uniform', 'distance'],
    'knn__p': [1, 2]
}

# Create GridSearchCV object for KNN with stratified cross-validation
grid_search_knn = GridSearchCV(
    estimator=knn_pipeline,
    param_grid=param_grid_knn,
    scoring='accuracy',
    cv=StratifiedKFold(n_splits=5, shuffle=True), # Stratified cross-validation
    n_jobs=-1,
    error_score='raise'
)

# Train the model on the resampled training data
grid_search_knn.fit(resampled_df.drop('churn', axis=1), resampled_df['churn'])

# Get the best estimator
```

```
# GET THE BEST ESTIMATOR
best_knn_model = grid_search_knn.best_estimator_

# Predict on the new_test_df using the best model
new_test_predictions_knn = best_knn_model.predict(new_test_df.drop('churn', axis=1))

# Calculate accuracy and F1-score on the new_test_df
accuracy_knn = accuracy_score(new_test_df['churn'], new_test_predictions_knn)
f1_knn = f1_score(new_test_df['churn'], new_test_predictions_knn)

print(f'Accuracy on test set (KNN): {accuracy_knn:.4f}')
print(f'F1-score on test set (KNN): {f1_knn:.4f}')
# Print best estimator's parameter values
print("Best Estimator's Parameters:")
print(grid_search_knn.best_params_)
```

Accuracy on test set (KNN): 0.8237  
 F1-score on test set (KNN): 0.5532  
 Best Estimator's Parameters:  
 {'knn\_n\_neighbors': 15, 'knn\_p': 1, 'knn\_weights': 'distance', 'pca\_n\_components': 9}

In [83]:


```
# Calculate accuracy, precision, recall, and F1-score on the new_test_df for KNN
accuracy_knn = accuracy_score(new_test_df['churn'], new_test_predictions_knn)
precision_knn = precision_score(new_test_df['churn'], new_test_predictions_knn)
recall_knn = recall_score(new_test_df['churn'], new_test_predictions_knn)
f1_knn = f1_score(new_test_df['churn'], new_test_predictions_knn)

print(f'Accuracy on test set (KNN): {accuracy_knn:.4f}')
print(f'Precision on test set (KNN): {precision_knn:.4f}')
print(f'Recall on test set (KNN): {recall_knn:.4f}')
print(f'F1-score on test set (KNN): {f1_knn:.4f}')
```

Accuracy on test set (KNN): 0.8237  
 Precision on test set (KNN): 0.4461  
 Recall on test set (KNN): 0.7280  
 F1-score on test set (KNN): 0.5532

In [84]:

```
# For KNN model
plot_confusion_matrix(best_knn_model, new_test_df)
```



## 9.4 RANDOM FORESTS

In [85]:

```
# Create a Random Forest Classifier
rf_classifier = RandomForestClassifier(random_state=42)

# Define the parameter grid for GridSearchCV
param_grid = {
    'n_estimators': [100, 200, 300], # Number of trees in the forest
    'max_depth': [None, 5, 10, 20], # Maximum depth of the tree
    'min_samples_split': [2, 5, 7], # Minimum number of samples required to split a node
    'min_samples_leaf': [1, 2, 4], # Minimum number of samples required at each leaf node
    'max_features': ['sqrt', 'log2'], # Number of features to consider for the best split
}
```

```
# Create GridSearchCV object for Random Forest
grid_search_rf = GridSearchCV(
    estimator=rf_classifier,
    param_grid=param_grid,
    scoring='accuracy', # Using accuracy as the scoring metric
    cv=3, # 3-fold cross-validation
    n_jobs=-1, # Use all available processors
)

# Train the model on the resampled training data (replace X_train and y_train with your data)
grid_search_rf.fit(resampled_df.drop('churn', axis=1), resampled_df['churn'])

# Get the best estimator
best_rf_model = grid_search_rf.best_estimator_

# Predict on the new_test_df using the best model
new_test_predictions_rf = best_rf_model.predict(new_test_df.drop('churn', axis=1))

# Calculate accuracy and F1-score on the new_test_df (replace y_true with true labels)
accuracy_rf = accuracy_score(new_test_df['churn'], new_test_predictions_rf)
f1_rf = f1_score(new_test_df['churn'], new_test_predictions_rf)

print(f'Accuracy on test set (Random Forest): {accuracy_rf:.4f}')
print(f'F1-score on test set (Random Forest): {f1_rf:.4f}')
```

Accuracy on test set (Random Forest): 0.9161  
F1-score on test set (Random Forest): 0.7266

In [86]:

```
# Calculate accuracy, precision, recall, and F1-score on the new_test_df for Random Forest
accuracy_rf = accuracy_score(new_test_df['churn'], new_test_predictions_rf)
precision_rf = precision_score(new_test_df['churn'], new_test_predictions_rf)
recall_rf = recall_score(new_test_df['churn'], new_test_predictions_rf)
f1_rf = f1_score(new_test_df['churn'], new_test_predictions_rf)

print(f'Accuracy on test set (Random Forest): {accuracy_rf:.4f}')
print(f'Precision on test set (Random Forest): {precision_rf:.4f}')
print(f'Recall on test set (Random Forest): {recall_rf:.4f}')
print(f'F1-score on test set (Random Forest): {f1_rf:.4f}')
```

Accuracy on test set (Random Forest): 0.9161  
Precision on test set (Random Forest): 0.7099  
Recall on test set (Random Forest): 0.7440  
F1-score on test set (Random Forest): 0.7266

## 9.5 XG-BOOST - CLASSIFIER

In [87]:

```
# Create an XGBoost classifier
xgb_classifier = XGBClassifier(random_state=42)

# Define the parameter grid for GridSearchCV
param_grid = {
    'n_estimators': [100, 200, 300], # Number of boosting rounds
    'max_depth': [3, 6, 9], # Maximum tree depth
    'learning_rate': [0.1, 0.01, 0.001], # Learning rate
    'subsample': [0.7, 0.8, 0.9], # Subsample ratio of the training instance
    'colsample_bytree': [0.7, 0.8, 0.9], # Subsample ratio of columns when constructing each tree
}

# Create GridSearchCV object for XGBoost
grid_search_xgb = GridSearchCV(
    estimator=xgb_classifier,
    param_grid=param_grid,
    scoring='accuracy', # Using accuracy as the scoring metric
    cv=3, # 3-fold cross-validation
    n_jobs=-1, # Use all available processors
)

# Train the model on the resampled training data (replace X_train and y_train with your data)
grid_search_xgb.fit(resampled_df.drop('churn', axis=1), resampled_df['churn'])

# Get the best estimator
best_xgb_model = grid_search_xgb.best_estimator_

# Predict on the new_test_df using the best model
new_test_predictions_xgb = best_xgb_model.predict(new_test_df.drop('churn', axis=1))

# Calculate accuracy and F1-score on the new_test_df (replace y_true with true labels)
accuracy_xgb = accuracy_score(new_test_df['churn'], new_test_predictions_xgb)
f1_xgb = f1_score(new_test_df['churn'], new_test_predictions_xgb)

print(f'Accuracy on test set (XGBoost): {accuracy_xgb:.4f}')
print(f'F1-score on test set (XGBoost): {f1_xgb:.4f}')
```

Accuracy on test set (XGBoost): 0.9245  
F1-score on test set (XGBoost): 0.7364

In [88]:

```
# Calculate accuracy, precision, recall, and F1 score on the new_test_df
```

```
# Calculate accuracy, precision, recall, and f1-score on the new_test_df
accuracy_xgb = accuracy_score(new_test_df['churn'], new_test_predictions_xgb)
precision_xgb = precision_score(new_test_df['churn'], new_test_predictions_xgb)
recall_xgb = recall_score(new_test_df['churn'], new_test_predictions_xgb)
f1_xgb = f1_score(new_test_df['churn'], new_test_predictions_xgb)

print(f'Accuracy on test set (XGBoost): {accuracy_xgb:.4f}')
print(f'Precision on test set (XGBoost): {precision_xgb:.4f}')
print(f'Recall on test set (XGBoost): {recall_xgb:.4f}')
print(f'F1-score on test set (XGBoost): {f1_xgb:.4f}')
```

Accuracy on test set (XGBoost): 0.9245  
Precision on test set (XGBoost): 0.7719  
Recall on test set (XGBoost): 0.7040  
F1-score on test set (XGBoost): 0.7364

In [ ]:

## 9.6 ENSEMBLING

In [89]:

```
# Get predictions from each base model
predictions_xgb = best_xgb_model.predict(new_test_df.drop('churn', axis=1))
predictions_rf = best_rf_model.predict(new_test_df.drop('churn', axis=1))
predictions_knn = best_knn_model.predict(new_test_df.drop('churn', axis=1))
predictions_tree = tree_pipeline.predict(new_test_df.drop('churn', axis=1))
predictions_logistic = best_model.predict(new_test_df.drop('churn', axis=1))

# Create a DataFrame with predictions and original features
predictions_df = new_test_df.drop('churn', axis=1).copy()
predictions_df['xgb_prediction'] = predictions_xgb
predictions_df['rf_prediction'] = predictions_rf
predictions_df['knn_prediction'] = predictions_knn
predictions_df['tree_prediction'] = predictions_tree
predictions_df['logistic_prediction'] = predictions_logistic

# Define the meta-Learner
meta_learner = LogisticRegression(random_state=42, max_iter=1000) # Increase max_iter

# Use predictions and original features to train the meta-Learner directly
meta_learner.fit(predictions_df, new_test_df['churn'])

# Scale the predictions and original features
scaler = StandardScaler()
scaled_predictions_df = scaler.fit_transform(predictions_df)

# Use scaled predictions and original features to train the meta-Learner
meta_learner.fit(scaled_predictions_df, new_test_df['churn'])

# Predict with the meta-Learner using the scaled combined dataset
final_predictions_scaled = meta_learner.predict(scaled_predictions_df)

# Evaluate the meta-Learner's performance with scaled data
accuracy_meta_scaled = accuracy_score(new_test_df['churn'], final_predictions_scaled)
f1_meta_scaled = f1_score(new_test_df['churn'], final_predictions_scaled)

print(f'Accuracy on test set (Meta-Learner with scaled data): {accuracy_meta_scaled:.4f}')
print(f'F1-score on test set (Meta-Learner with scaled data): {f1_meta_scaled:.4f}')
```

Accuracy on test set (Meta-Learner with scaled data): 0.9317  
F1-score on test set (Meta-Learner with scaled data): 0.7554

In [90]:

```
# Calculate accuracy, precision, recall, and F1-score on the new_test_df for the meta-Learner with scaled data
accuracy_meta_scaled = accuracy_score(new_test_df['churn'], final_predictions_scaled)
precision_meta_scaled = precision_score(new_test_df['churn'], final_predictions_scaled)
recall_meta_scaled = recall_score(new_test_df['churn'], final_predictions_scaled)
f1_meta_scaled = f1_score(new_test_df['churn'], final_predictions_scaled)

print(f'Accuracy on test set (Meta-Learner with scaled data): {accuracy_meta_scaled:.4f}')
print(f'Precision on test set (Meta-Learner with scaled data): {precision_meta_scaled:.4f}')
print(f'Recall on test set (Meta-Learner with scaled data): {recall_meta_scaled:.4f}')
print(f'F1-score on test set (Meta-Learner with scaled data): {f1_meta_scaled:.4f}')
```

Accuracy on test set (Meta-Learner with scaled data): 0.9317  
Precision on test set (Meta-Learner with scaled data): 0.8148  
Recall on test set (Meta-Learner with scaled data): 0.7040  
F1-score on test set (Meta-Learner with scaled data): 0.7554

## 9.7 MODELS AND THEIR PERFORMANCES

In [91]:

```
# Define model names
model_names = ['XGBoost', 'Random Forest', 'KNN', 'Decision Tree', 'Logistic Regression', 'Stacked Model ']

# Define accuracy scores
accuracy_scores = [accuracy_xgb, accuracy_rf, accuracy_knn, accuracy_tree, accuracy, accuracy_meta_scaled] # Replace with accuracy values

# Create a DataFrame to store model names and accuracy scores
```

```
accuracy_df = pd.DataFrame({
    'Model': model_names,
    'Accuracy Score': accuracy_scores
})

# Display the accuracy scores as a table
print(accuracy_df)
```

	Model	Accuracy Score
0	XGBoost	0.924460
1	Random Forest	0.916067
2	KNN	0.823741
3	Decision Tree	0.769784
4	Logistic Regression	0.784173
5	Stacked Model	0.931655

In [92]:

```
combined_scores = {
    'XGBoost': {
        'Accuracy': accuracy_xgb,
        'Precision': precision_xgb,
        'Recall': recall_xgb,
        'F1-score': f1_xgb
    },
    'Random Forest': {
        'Accuracy': accuracy_rf,
        'Precision': precision_rf,
        'Recall': recall_rf,
        'F1-score': f1_rf
    },
    'KNN': {
        'Accuracy': accuracy_knn,
        'Precision': precision_knn,
        'Recall': recall_knn,
        'F1-score': f1_knn
    },
    'Decision Tree': {
        'Accuracy': accuracy_tree,
        'Precision': precision_tree,
        'Recall': recall_tree,
        'F1-score': f1_tree
    },
    'Meta_learner_scaled_scores': {
        'Accuracy': accuracy_meta_scaled,
        'Precision': precision_meta_scaled,
        'Recall': recall_meta_scaled,
        'F1-score': f1_meta_scaled
    }
}
```

In [93]:

```
# Create Lists to store model names and corresponding scores
models = []
accuracy_scores = []
precision_scores = []
recall_scores = []
f1_scores = []

# Loop through the combined_scores dictionary to extract scores for each model
for model, scores in combined_scores.items():
    models.append(model)
    accuracy_scores.append(scores['Accuracy'])
    precision_scores.append(scores['Precision'])
    recall_scores.append(scores['Recall'])
    f1_scores.append(scores['F1-score'])

# Create a DataFrame using pandas
scores_table = pd.DataFrame({
    'Model': models,
    'Accuracy Score': accuracy_scores,
    'Precision Score': precision_scores,
    'Recall Score': recall_scores,
    'F1 Score': f1_scores
})

# Display the table
print(scores_table)
```

	Model	Accuracy Score	Precision Score	Recall Score	\
0	XGBoost	0.924460	0.771930	0.704	
1	Random Forest	0.916067	0.709924	0.744	
2	KNN	0.823741	0.446078	0.728	
3	Decision Tree	0.769784	0.356223	0.664	
4	Meta_learner_scaled_scores	0.931655	0.814815	0.704	
	F1 Score				
0		0.736402			
1		0.726562			
2		0.553191			
3		0.463687			
4		0.755365			

```
In [94]: from tabulate import tabulate

# Create lists to store the table data
table_data = []
headers = ['Model', 'Accuracy Score', 'Precision Score', 'Recall Score', 'F1 Score']

# Populate the table data
for model, scores in combined_scores.items():
    table_data.append([model, scores['Accuracy'], scores['Precision'], scores['Recall'], scores['F1-score']])

# Use tabulate to create a pretty table
pretty_table = tabulate(table_data, headers=headers, tablefmt="pretty")

# Display the pretty table
print(pretty_table)
```

Model	Accuracy Score	Precision Score	Recall Score	F1 Score
XGBoost	0.9244604316546763	0.7719298245614035	0.704	0.7364016736401674
Random Forest	0.9160671462829736	0.7099236641221374	0.744	0.7265625
KNN	0.8237410071942446	0.44607843137254904	0.728	0.5531914893617021
Decision Tree	0.7697841726618705	0.3562231759656652	0.664	0.46368715083798884
Meta_learner_scaled_scores	0.9316546762589928	0.8148148148148148	0.704	0.7553648068669528

In [ ]:

## 10 DEPLOYMENT

```
In [95]: # Save the model to a file
joblib.dump(meta_learner, 'meta_learner.joblib')
```

Out[95]: ['meta\_learner.joblib']

In [ ]:

In [ ]: